

CNS HW2

B03902082 資工四 江懿友

1. Super Cookie

1.1 HSTS

HSTS is a mechanism to protect against "protocol downgrade attack" and "cookie hijacking attack". It works by forcing the client (web browser) to communicate with the server only over https and never via http. Upon a client's first visit to a website, a HSTS header will be sent over https, instructing the client that "any future connection to this domain of http should be automatically converted to https connection" and "supply a error message to the user if https connection cannot be established when visiting this domain".

1.2 Super Cookie via HSTS

The tracker can assign for each user a unique ID, say a 32-bit integer. Then on the first visit of the user, the tracker sets the HSTS policy of his 32 subdomains according to the user ID, and the 32 subdomains are accessed by the user by 32 images embedded in the webpage. Thus 32 1-bit super cookies can be stored on the user's browser. When the user visits the site again, he would access the 32 images of different subdomains, and access them with http or https according to the users HSTS record. So by checking if his 32 subdomains are accessed via http or https, the tracker can recover the user's unique ID.

Reference: <https://thehackernews.com/2018/03/hsts-supercookie-tracking.html>

1.3 Mitigations

Only allow setting HSTS for the visiting hostname. For example if I'm visiting "www.example.org" then only the HSTS setting of "www.example.org" should be set but not "sub.www.example.org" nor "sub.example.org". This way the tracker cannot abuse HSTS for tracking by browser automatically loading resources from tracker's subdomains.

2. BGP

2.1 Prefix Hijacking

It is likely that the attacker is performing "BGP prefix hijacking".

The attacker announces itself to be the owner of "10.10.220.0/23" and "10.10.222.0/23". Since BGP routing prefers the longest matching prefix, the other ASs would route destinations of "10.10.220.0/22" to AS999 instead of AS1000 because AS999 has a longer match than AS1000 ($23 > 22$).

2.2.a Spurious BGP Update

UPDATE {"10.10.220.0/23", {AS999, AS2, AS1, AS1000}}

UPDATE {"10.10.222.0/23", {AS999, AS2, AS1, AS1000}}

2.2.b Explain 2.2.a

For AS3, AS4 and AS5, they would route destination "10.10.220.0/22" to AS999 because the route to AS999 have a longer matching prefix.

For AS1 and AS2, they won't accept AS999's spurious BGP update according to the "loop prevention" rule.

2.2.c

Advantage: AS999 doesn't have to pretend to be the owner of "10.10.220.0/22" like in 2.1. So if AS3, AS4 or AS5 were to verify the owner of "10.10.220.0/23" and "10.10.222.0/23" then they would find out that AS1000 indeed owns "10.10.220.0/23" and "10.10.222.0/23".

Disadvantage: The attacker cannot tamper with packets originating from AS1 and AS2.

3. PIN Authentication

The server S can be used as an oracle that answers "Does X equals PIN1?" where X is the guess of PIN1.

First let's guess PIN1 to be X. At step (3) we send $\text{HMAC_K1}(\text{RC1} \parallel X)$ to server. If X does not equal PIN1 then the server would close the connection after step (5). If X does equal PIN1 then the server would proceed to step (6) and send some message back, thus indicating I made a right guess. The total combination of PIN1 is $10^4 = 10000$ so I can use this oracle to brute force PIN1 in feasible time.

After obtaining PIN1 a similar process can be used to search PIN2. Only this time I use "whether or not server close connection at step (8)" as a indicator of whether or not I made a right guess. Thus I can recover the whole PIN in feasible time.

The vulnerability of the protocol is that it aborts immediately when a verification fails, and this leaks information to the attacker. Another vulnerability is that the PIN is split in half and the two halves are used separately. This reduced the effective key length from 8 digits to 4 digits and made my brute force attack possible.

4. Can't Beat CBC

4.1

Flag = BALS{IV=KEY=GG}

Suppose

$m = \text{'QQ Homework is too hard, how2decrypt QQ'}$

$c =$

$\text{'296e12d608ad04bd3a10b71b9eef4bb6ae1d697d1495595a5f5b98e409d7a7c437f24e69feb250b347db0877a40085a9'}$

where m is the plain text and c is cipher text

Then I partition the plain and cipher text into three blocks

$m = M1 \parallel M2 \parallel M3$

$c = C1 \parallel C2 \parallel C3$

Then I ask the server to decrypt the message $M1 \parallel C1 \parallel C2 \parallel C3$.

So when CBC operation mode tries to decrypt this message it does:

$D(M1) \oplus IV = P1$

$D(C1) \oplus M1 = P2$

$D(C2) \oplus C1 = P3$

$D(C3) \oplus C2 = P4$

$P1$ is some random value that we don't care.

Because $D(C1) = IV \oplus M1$, so $P2 = IV$ which is the flag.

And $D(C3) = C2 \oplus M3$ so $P4 = M3$. Thus the decrypted message have proper padding.

4.2

Flag =

BALS{1T_15_V3RY_FUN_T0_533_TH3_FL4G_4PP34R_0N3_BY_0N3_R1GHT_XD}

This problem can be solved by CBC padding oracle attack. Searching for the plaintext byte by byte is time consuming, so the connection to server may timeout, thus I have to reconnect to the server periodically.

4.3

Flag =

BALSN{N33D_B3TT3R_P4DD1NG_M3TH0D_T0_PR3V3NT_P00DL3_4TT4CK_SS
Lv3}

The vulnerability is that the MAC is checked after the padding has been removed, so I can alter the ciphertext without being detected by MAC checking by carefully choosing the padding length so that when the server un pads it removes my altered message.

My attack roughly works like:

1. Ask server to encrypt FLAG, so I get $(C1 \parallel C2 \parallel \text{MAC}(C1 \parallel C2) \parallel \text{padding})$, where $C1 \parallel C2$ is the encrypted FLAG.
2. Ask server to decrypt $(C1 \parallel C2 \parallel \text{MAC}(C1 \parallel C2) \parallel \text{padding} \parallel C1' \parallel C2)$, where $C1'$ is $C1$ block modified using a way similar to CBC padding oracle attack. I try to guess the last byte of $C1$'s plaintext, if my guess is correct then the "padding $\parallel C1' \parallel C2$ " can be correctly removed when server un pads and the MAC checking can pass.
3. After I find the last byte of $C1$'s plaintext, I ask server to encrypt 'a' \parallel FLAG, so I get $(\text{'a'} \parallel C1 \parallel C2 \parallel \text{MAC}(\text{'a'} \parallel C1 \parallel C2) \parallel \text{padding})$. The ciphertext is shifted by one byte, so when I use the same procedure to recover the last byte of $C1$'s plaintext I actually get the second to last byte of the plaintext.

So by repeating this procedure I can eventually recover the whole FLAG.

5. Man-in-the-Middle Attack

Flag = BALSN{Wow_you_are_really_in_the_middle}

This problem is vulnerable to reflection attack.

Let pwd1 , pwd2 , pwd3 be the three unknown parts of the password.

First I initiate two connections with the server, and I receive $B11 = g^{1^b11}$ and $B21 = g^{1^b21}$ from the two connections, where $b11$, $b21$ and g is unknown by me.

Then I send $B21$ to connection 1 and $B11$ to connection 2. This way connection 1 and 2 would both derive the secret $B21^{b11} = g^{1^{(b11 \cdot b21)}} = B11^{b21}$.

The same goes for the second key exchange, I forward $B12 = g^{2^b12}$ to connection 2 and $B22 = g^{2^b22}$ to connection 1. So now connection 1 and connection 2 both derives the same secret.

Then I will receive $B_{13} = g_3^{b_{13}}$, $B_{23} = g_3^{b_{23}}$ from the two connections. I then try to guess the value of g_3 and send $A = g_3^1$ to both connections. So connection 1 derives the secret $g_3^{b_{13}}$ and connection 2 derives $g_3^{b_{23}}$.

Then I'll receive two encrypted messages. The messages are the FLAG xor masked with the established shared secrets. Because connection 1 and connection 2's shared secret only differ at the last part (the part derived from g_3), so when I xor the two messages together I get $(A^{b_{13}} \text{ xor } A^{b_{23}})$. If my guess of g_3 is correct then this value will equal to $(B_{13} \text{ xor } B_{23})$ else there is a high chance that $(A^{b_{13}} \text{ xor } A^{b_{23}}) \neq (B_{13} \text{ xor } B_{23})$. So now I have a way to verify my guess, I can search the value of g_3 by trying pwd_3 from 1 to 20, as g_3 is calculated from pwd_3 .

Then I can use the same procedure to search g_2 and g_1 . After finding all of g_1 , g_2 and g_3 , I can recover the encrypted FLAG.

6. Cloudburst

Flag = BALS{what_a_CloudPiercer}

Origin IP: 140.112.91.250

Inspired by the Cloudpiercer paper ([link](#)), I scan the csie department IP region (according to this table <http://ccnet.ntu.edu.tw/NTU/IP/IPserver.html>) and dumped all of the hosts' SSL certificate. My assumption is that knowing the origin host is listening on port 443, it is likely that origin host's https service uses a certificate whose CN includes "the-real-reason-to-not-use-sigkill.csie.org".

This assumption is proved to be correct as I dumped a SSL certificate from the IP 140.112.91.250 that shows:

subject=/CN=the-real-reason-to-not-use-sigkill.csie.org

issuer=/C=US/O=Let's Encrypt/CN=Let's Encrypt Authority X3

7. One-time Wallet

Flag = BALS{R3tir3_4t_tw3nty}

The vulnerability is that the address / password generator uses `random.getrandbits()` to generate pseudo random bit sequences, however the PRNG is actually predictable.

The PRNG used by `random.getrandbits()` is Mersenne twister. Mersenne twister's internal state can be determined by 624 consecutive outputs. In other words, knowing the first 624 outputs of a Mersenne twister, one can predict the 625-th and the rest of output.

I used this tool to predict Mersenne twister:
pip3 install mersenne-twister-predictor

8. TLS Certificate

Flag = BALS{t1s_ChAiN_0f_7ru5t}

This problem is rather straightforward. I read OpenSSL's documentation to learn how to sign a X509 certificate.

First I downloaded csie_cert from <https://www.csie.ntu.edu.tw/> .

Then I craft a new X509 certificate, and I copy csie_cert's subject and publickey into my certificate. I also set my certificate's serial number, not_before and not_after to some arbitrary value. At last I set my certificate's issuer to BALS rootCA's subject. Finally I sign my certificate with the provided RSA private key and send the signed CA to server.